



A
Project Report
on
**Identification of Skin Abnormalities via Deep Learning:
An Optimized ResNet50 Approach**
submitted as partial fulfillment for the award of
**BACHELOR OF TECHNOLOGY
DEGREE**

SESSION 2025-26

in
Computer Science and Engineering

By

Ankur Kumar (2200290100026)

Anshul Rai (2200290100031)

Anmol Kumar (2200290100027)

Abhishek Yadav (2200290100008)

Under the supervision of

Mr. Pushpendra Kumar

KIET Group of Institutions, Ghaziabad

Affiliated to

**Dr. A.P.J. Abdul Kalam Technical University,
Lucknow**
(Formerly UPTU)

May, 2026

DECLARATION

We hereby declare that this submission is our own work and that, to the best of our knowledge and belief, it contains no material previously published or written by another person nor material which to a substantial extent has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text.

Roll No.	Name	Signature	Date
2200290100026	ANKUR KUMAR		
2200290100031	ANSHUL RAI		
2200290100027	ANMOL KUMAR		
2200290100008	ABHISHEK YADAV		

CERTIFICATE

This is to certify that Project Report entitled “Identification of Skin Abnormalities via Deep Learning: An Optimized ResNet50 Approach” which is submitted by Ankur Kumar, Anshul Rai, Anmol Kumar and Abhishek Yadav in partial fulfillment of the requirement for the award of degree B. Tech. in Department of Computer Science and Engineering of Dr. A.P.J. Abdul Kalam Technical University, Lucknow is a record of the candidates own work carried out by them under my supervision. The matter embodied in this report is original and has not been submitted for the award of any other degree.

.

Pushpendra Kumar

(Assistant Professor)

Dr. Vineet Sharma

(Dean CSE)

Date:

ACKNOWLEDGEMENT

It gives us a great sense of pleasure to present the report of the B. Tech Project undertaken during B. Tech. Final Year. We owe a special gratitude to Mr. Pushpendra Kumar, Department of Computer Science and Engineering, KIET, Ghaziabad, for his constant support and guidance throughout the course of our work. His sincerity, thoroughness and perseverance have been a constant source of inspiration for us. It is only his cognizant efforts that our endeavors have seen light of the day.

We also take the opportunity to acknowledge the contribution of Dr. Vineet Sharma, Dean of the Department of Computer Science and Engineering, KIET, Ghaziabad, for his full support and assistance during the development of the project. We also do not like to miss the opportunity to acknowledge the contribution of all the faculty members of the department for their kind assistance and cooperation during the development of our project. . Last but not the least, we acknowledge our friends for their contribution in the completion of the project.

Roll No.	Name	Signature	Date
2200290100026	ANKUR KUMAR		
2200290100031	ANSHUL RAI		
2200290100027	ANMOL KUMAR		
2200290100008	ABHISHEK YADAV		

ABSTRACT

Skin diseases are among the most common health problems worldwide and affect millions of individuals across different age groups and geographical locations. Dermatological conditions such as acne, eczema, rosacea, herpes, and fungal infections can significantly impact a person's physical and psychological well-being. Early detection and accurate diagnosis of these diseases are essential for effective treatment and prevention of complications. In many rural and underdeveloped regions, the lack of specialized dermatologists makes early diagnosis difficult.

With the advancement of Artificial Intelligence and Deep Learning technologies, automated diagnostic systems have gained significant attention in the medical field. Deep learning models, particularly Convolutional Neural Networks (CNNs), have demonstrated exceptional performance in analyzing medical images and identifying complex patterns that may not be easily visible to the human eye.

This project proposes a deep learning based framework for the identification of skin abnormalities using an optimized ResNet50 architecture. The proposed system follows a structured pipeline that includes image preprocessing, data augmentation, model training, and evaluation. A two-stage training approach is implemented in which the classification head is first trained while freezing the base layers of the network, followed by fine-tuning of deeper layers to capture domain-specific features.

The performance of the system is evaluated using metrics such as accuracy, precision, recall, and F1-score. Experimental results demonstrate that the optimized ResNet50 model achieves high classification accuracy and is capable of effectively identifying various skin abnormalities.

TABLE OF CONTENTS	Page No.
DECLARATION.....	ii
CERTIFICATE.....	iii
ACKNOWLEDGEMENTS.....	iv
ABSTRACT.....	v
LIST OF FIGURES.....	viii
LIST OF TABLES.....	ix
LIST OF ABBREVIATIONS.....	x
CHAPTER 1 (INTRODUCTION)	1
1.1. Introduction.....	1
1.2. Project Description.....	2
1.3. Problem Statement.....	4
1.4. Objectives.....	5
1.5. Scope of the Project.....	6
CHAPTER 2 (LITERATURE REVIEW)	8
2.1. Overview Of Skin Disease Detection Systems.....	8
2.2. Research Gap.....	11
2.3 Motivation for the Proposed System.....	13
CHAPTER 3 (PROPOSED METHODOLOGY)	16
3.1. System Overview.....	16
3.2. Data Collection	17
3.3. Data Preprocessing.....	18
3.4. Data Augmentation.....	20
3.5. Model Architecture.....	21
3.6 Training Strategy.....	23
3.7. Evaluation Metrics.....	25

CHAPTER 4 (RESULTS AND DISCUSSION)	27
4.1. Data Analysis.....	27
4.2. Model Training Perfomance.....	28
4.3. Confusion Matrix Analysis.....	30
4.4. Discussion.....	33
 CHAPTER 5 (CONCLUSIONS AND FUTURE SCOPE)	 36
5.1. Conclusion.....	36
5.2. Future Scope.....	39
 REFERENCES.....	 42
 APPENDEX 1	 44

LIST OF FIGURES

Figure No.	Description	Page No.
2.1	Optimized ResNet50 Diagnostic Pipeline for skin Disease Classification, shows how the raw data is used to recognise the pattern to predict the disease.	11
3.3	Data Augementated process for skin disease lesion robustness.	19
4.2	This figure shows the accuracy and loss graph by using 50 epoch on training data and testing data.	28
4.3	Confusion Matrix of Skin Diseases Predicted.	32

LIST OF TABLES

Table. No.	Description	Page No.
3.7	Comparative analysis of deep learning architectures for skin disease classification	24
4.4	Overall performance summary of the proposed ResNet50 model	35

LIST OF ABBREVIATIONS

Abbreviation	Full Form
AI	Artificial Intelligence
CNN	Convolutional Neural Network
DL	Deep Learning
CAD	Computer Aided Diagnosis
GPU	Graphics Processing Unit
TP	True Positive
FP	False Positive
FN	False Negative
TN	True Negative
F1	Harmonic Mean of Precision and Recall
PIL	Python Imaging Library
GAP	Global Average Pooling
VGG16	Visual Geometry Group 16-layer model
YOLOv5	You Only Look Once
ResNet50	Residual Network with 50 layers
ReLU	Rectified Linear Unit

CHAPTER 1

INTRODUCTION

1.1 Introduction

Skin diseases represent a profound, ubiquitous, and rapidly escalating burden on global public health systems. Millions of people get affected by skin diseases every year, and these conditions span across all age groups, geographic locations, and socioeconomic boundaries. The spectrum of these dermatological conditions is incredibly vast and highly heterogeneous, ranging from common, highly treatable ailments to severe, life-threatening malignancies. Conditions like acne, eczema, psoriasis, and various forms of skin cancer not only affect the physical health and longevity of people but also inflict significant psychological distress. Patients frequently experience social stigma, severe anxiety, and a heavily reduced quality of life due to the visible nature of these ailments. Furthermore, the economic burden associated with chronic skin conditions—including the cost of prolonged treatments, frequent clinical visits, and lost occupational productivity—is staggering.

Despite the prevalence of these conditions, they pose a formidable challenge in diagnosis due to diverse clinical presentations and massive overlaps in visual appearance. The sheer visual similarity between a benign lesion (such as a harmless melanocytic nevus or a seborrheic keratosis) and a malignant tumor (such as an early-stage melanoma or basal cell carcinoma) often complicates the initial screening process. This diagnostic overlap requires a highly trained clinical eye to discern the microscopic, highly subtle differences in texture, color distribution, and border irregularity.

Compounding this inherent biological complexity is a significant systemic issue in global healthcare infrastructure. Access to specialist dermatologists is severely limited in many regions across the globe, which inevitably leads to incorrect or dangerously delayed diagnoses. In rural, remote, or developing areas, patients frequently rely on general medical practitioners who may lack the specialized, niche training required to accurately differentiate

between visually confounding skin abnormalities. Early accurate diagnosis is critically important, especially for life-risking diseases like skin cancer. A delayed diagnosis in these critical cases allows the disease to metastasize and spread beyond the primary site, severely reducing the efficacy of subsequent surgical or oncological treatments and exponentially increasing mortality rates.

To address these widening gaps in healthcare accessibility and to mitigate the risks of human error, the medical field has increasingly turned to advanced computational solutions. An automated image-based diagnosis model can help in assisting medical professionals and provide immediate, objective preliminary screening. Through the rapid evolution of Deep Learning, particularly the advent of Convolutional Neural Networks (CNNs), the landscape of medical image analysis has changed significantly. These advanced neural architectures bypass the need for manual, subjective feature extraction—which was a major, crippling limitation of traditional, classical machine learning paradigms. Instead, CNNs automatically learn complex, hierarchical visual representations directly from raw pixel data.

Various complex network architectures have been rigorously tested in this medical domain. Deep learning models like ResNet, VGGNet, and InceptionNet have demonstrated remarkable performance in skin lesion classification, sometimes even surpassing human expectations in controlled environments. Among these architectures, ResNet50 has emerged as a particularly powerful and robust tool. Therefore, the medical computing field heavily relies on models like ResNet50, which is immensely popular because of its optimal balance between structural depth and computational efficiency. The innovative use of residual "skip connections" in ResNet50 allows the network to maintain a massive depth of 50 layers without suffering from the "vanishing gradient problem"—a mathematical phenomenon where learning stalls in deep networks. This unique architecture enables it to extract the highly complex, abstract visual patterns necessary for accurate dermatological classification.

1.2 Project Description

This research project, titled "Identification of Skin Abnormalities via Deep Learning: An Optimized ResNet50 Approach," focuses on the rigorous development, architectural optimization, and statistical evaluation of a robust deep learning framework tailored

specifically for dermatological analysis. This paper presents an optimized ResNet50 model specifically engineered for multi-class skin disease classification.

The core of the proposed diagnostic system leverages the concept of transfer learning. Transfer learning is a highly efficient machine learning paradigm wherein a complex neural network, originally pre-trained on a massive, general-purpose image dataset, is meticulously fine-tuned to recognize the specific, nuanced visual signatures of a new domain—in this case, dermatological conditions. Specifically, we have utilized and fine-tuned a ResNet50 model that was previously pretrained on the ImageNet database, leveraging its pre-existing capacity to recognize fundamental visual components like edges, curves, and color gradients.

The model was trained on a comprehensive dataset encompassing 22 different skin diseases. Because raw medical data cannot be fed directly into a deep neural network without causing severe computational instability, this project constructs a highly structured, multi-stage data processing pipeline. The main goal of this project is to drastically improve the predictive accuracy and real-world generalization of the model through a synergistic combination of transfer learning, aggressive data augmentation, and advanced regularization techniques.

To optimize the architecture for this specific multi-class medical task, a custom training and fine-tuning strategy was implemented. The strategy involves freezing the earlier layers of the network to preserve foundational visual knowledge, while selectively unfreezing later layers of the ResNet50 model. This selective unfreezing helps in the precise adaptation to skin-specific features while preserving the generic visual knowledge acquired from ImageNet.

Furthermore, to combat the inherent limitations of training deep networks on restricted medical datasets, we have deeply optimized the system. We employed extensive data augmentation techniques, including random horizontal and vertical flips, rotational shifts, dynamic zooming, and contrast adjustments. To interpret the complex feature maps generated by the ResNet50 backbone, a custom multi-layer classifier was appended to the network's apex. A learning rate scheduler was also implemented to dynamically adjust the learning step size, allowing the model to converge smoothly toward the optimal solution. Ultimately, this automated image-based classification model is designed to provide early diagnostic indicators, thereby lowering the risk of severe damage to the patient's body through timely medical intervention.

1.3 Problem Statement

While deep learning models hold immense, transformative promise for medical image analysis, transitioning these theoretical architectures from pristine laboratory frameworks to chaotic, real-world clinical applications reveals significant systemic challenges. We cannot directly apply such off-the-shelf models in real-world clinical problems without encountering severe analytical failures. The accurate, reliable, and unbiased automated classification of skin diseases remains a formidable computational challenge due to several compounding factors intrinsically related to the nature of medical data.

First and foremost, medical datasets are notoriously difficult to compile. They are often small, highly imbalanced, and contain massive variability in lighting, camera scale, skin tone, and other environmental imaging conditions. Unlike datasets used for general object recognition (which contain millions of perfectly balanced and well-lit images of everyday items), dermatological datasets are heavily constrained by strict patient privacy laws, ethical data collection hurdles, and the immense, expensive labor required for expert clinical annotation.

The single most critical problem addressed in this research is the highly unequal distribution of data across different disease categories—a phenomenon known as class imbalance. The financial and logistical cost of acquiring a labeled dataset for rare or aggressive diseases inevitably leads to severe imbalances in multi-class datasets. In real-world clinical settings, certain benign conditions (like simple nevi or common eczema) appear with vastly higher frequency than rare pathological infections or specific carcinomas. This numerical imbalance can make the neural network heavily biased towards the majority classes, resulting in poor generalization and highly unreliable predictions for uncommon, yet potentially fatal, minority diseases.

When a deep neural network is trained on highly imbalanced data, it falls victim to the "accuracy paradox." The model naturally gravitates toward predicting the most frequent classes because doing so superficially minimizes its mathematical loss function. For example, if a dataset is 80% benign, a model that simply guesses "benign" every single time will achieve 80% overall accuracy, yet it completely fails its medical purpose by missing every single critical minority case. To address such debilitating problems, researchers must explore various sophisticated optimization strategies, including transfer learning with pre-

trained weights, algorithmic data augmentation, specific architectural modifications, and stringent regularization techniques. Therefore, the central problem this project tackles is how to systematically optimize a ResNet50 architecture to overcome these severe dataset imbalances and variations in imaging conditions, ultimately producing a reliable, unbiased multi-class diagnostic tool.

1.4 Objectives

To effectively confront and mathematically neutralize the challenges outlined in the problem statement, this research is driven by a precise set of academic, structural, and technical objectives. These primary objectives are meticulously designed to systematically evaluate how structural and algorithmic optimizations can enhance deep learning performance in constrained, imbalanced medical domains:

- **Develop a Comprehensive Multi-Class Classification Framework:** To design, build, and implement an optimized ResNet50 deep learning model capable of accurately classifying a wide and complex spectrum of dermatological conditions. This specifically targets a diverse dataset comprising 22 distinct skin disease categories, moving beyond simplistic binary (benign vs. malignant) classification models.
- **Implement and Fine-Tune Advanced Transfer Learning Paradigms:** To systematically adapt the ResNet50 architecture by fine-tuning a model pre-trained on the ImageNet repository. This objective explicitly involves selectively freezing the initial 140 layers of the network to retain robust low-level feature extraction capabilities (such as edge and color detection), while actively training the remaining deeper layers to capture abstract, domain-specific dermatological textures and patterns.
- **Enhance Model Generalization via Algorithmic Optimization:** To drastically improve the diagnostic accuracy and real-world generalization of the model through a hybrid combination of transfer learning, spatial data augmentation, and regularization techniques. This objective explicitly includes applying extensive, randomized data augmentation techniques like random flip, rotation, zoom, and

contrast modification to artificially expand the dataset. Furthermore, it involves designing a custom multi-layer dense classifier and utilizing a dynamic learning rate scheduler to optimize the gradient descent process.

- **Stabilize Neural Training and Mitigate Catastrophic Overfitting:** To actively combat the well-documented overfitting tendencies of excessively deep networks when exposed to small medical datasets. This objective ensures that the training pipeline is mathematically stabilized using regularization pipelines, most notably the implementation of Early Stopping callbacks to halt training when validation performance begins to degrade.
- **Analyze, Expose, and Address Severe Class Imbalance:** A crucial objective is to rigorously evaluate the model's performance using metrics beyond raw accuracy. By utilizing macro-averaged and weighted precision, recall, and F1-scores, the objective is to expose the algorithmic challenge of class imbalance in medical image classification and meticulously highlight the necessity for more sophisticated future strategies, such as class-weighted loss functions or artificial data generation.

1.5 Scope of the Project

Clearly defining the boundaries of the research is essential to establish the operational parameters and clinical applicability of the proposed system. The scope of this project is explicitly defined by the specific technological methodologies applied, the software environment utilized, and the strict boundaries of the medical dataset utilized for training and evaluation.

The project focuses exclusively on the software engineering, algorithmic optimization, and detailed statistical analysis of a ResNet50 Convolutional Neural Network. The hardware and software scope includes the utilization of the Python programming language, specifically leveraging the TensorFlow and Keras deep learning frameworks for model construction and execution.

The data scope is strictly limited to a publicly available, pre-compiled skin disease dataset sourced directly from KaggleHub (specifically the creatorai/skin-disease repository), which

contains clinical RGB images categorized into 22 distinct skin conditions. The images within this dataset are programmatically resized and standardized into 224 x 224 pixel dimensions to strictly match the rigid input tensor requirements of the ResNet50 base architecture. The operational scope of model training involves systematically dividing the complete dataset into an 80/10/10 split, where 80% of the images are allocated for the active training phase, 10% are reserved for validation (hyperparameter tuning and Early Stopping checks), and the final 10% are sequestered for final, unbiased testing.

Furthermore, the evaluation scope of the project extends far beyond the reporting of simple categorical accuracy. The project encompasses a rigorous quantitative and qualitative evaluation utilizing standard classification metrics, including macro-averaged and weighted-averaged precision, recall, and F1-scores. This is conducted alongside a highly detailed macroscopic analysis of the generated confusion matrix to explicitly identify and document majority-class biases.

The project acknowledges certain boundary limitations within its scope. It does not involve the collection of novel patient data, nor does it encompass any live clinical trials or physical deployment in a hospital setting. Furthermore, while extensive spatial data augmentation was utilized to artificially inflate the data volume, the dataset utilized was not perfectly balanced, and no explicit algorithmic class-balancing mechanisms (such as Synthetic Minority Over-sampling Technique or Focal Loss) were implemented in this current iteration. Consequently, the scope of future work will focus heavily on addressing this class imbalance and transitioning the model toward edge-device clinical operations. The current iteration serves as a vital foundational study, unequivocally demonstrating the absolute necessity of advanced optimization when applying deep residual networks to highly imbalanced, multi-class dermatological data.

CHAPTER 2

LITERATURE REVIEW

2.1 Overview of Skin Disease Detection Systems

The intersection of computer science and dermatology has witnessed exponential growth over the past two decades. As global healthcare systems grapple with increasing patient loads and a distinct shortage of specialized medical professionals, the academic community has aggressively pursued computational methods to automate diagnostic workflows. This chapter provides a comprehensive, critical examination of the existing literature surrounding automated skin disease classification. It traces the evolutionary timeline from rudimentary image processing algorithms to sophisticated deep learning architectures. By critically analyzing contemporary methodologies, this review isolates the persistent technological and statistical shortcomings in the field—most notably the pervasive challenges of dataset imbalance and model overfitting. Ultimately, this review establishes the theoretical foundation and the driving motivation for deploying an optimized, transfer-learning-based ResNet50 architecture to tackle complex, multi-class dermatological diagnostics.

The historical progression of computational dermatology can be distinctly divided into two major technological eras: the classical machine learning epoch, which relied on explicit human programming, and the contemporary deep learning era, characterized by autonomous feature extraction. The earliest attempts to automate skin lesion analysis were highly constrained binary systems, designed simply to output a "healthy" or "diseased" verdict. However, as computational power scaled, clinical demands expanded, leading to the current focus on highly complex, multi-class classification matrices that attempt to mirror a human dermatologist's differential diagnosis.

Before the proliferation of high-capacity graphical processing units (GPUs), automated diagnostic systems were strictly bound by the limitations of classical digital image processing. In these foundational models, the system could not "see" or "learn"

independently. Instead, human engineers, working in tandem with medical experts, were required to manually translate clinical heuristics into rigid mathematical formulas.

A primary example of this was the algorithmic digitization of the widely used ABCD rule (Asymmetry, Border irregularity, Color variegation, and Diameter). Researchers developed complex contour-tracing algorithms to mathematically measure the irregularity of a lesion's border. To analyze texture, computational models frequently employed algorithms like the Gray-Level Co-occurrence Matrix (GLCM), which statistically mapped the spatial frequencies of pixel intensities to detect the rough, scaly surfaces indicative of conditions like psoriasis or squamous cell carcinoma. Color analysis was similarly manual, involving the extraction of multi-dimensional color histograms across various color spaces (such as RGB, HSV, or CIELAB) to quantify the presence of dangerous blue-white veils or atypical pigment networks.

Once these specific geometric, textural, and photometric features were meticulously calculated, they were aggregated into one-dimensional feature vectors. These vectors were subsequently fed into traditional machine learning classifiers to generate a diagnosis. The dominant algorithms of this era included Support Vector Machines (SVMs), which sought to calculate an optimal mathematical hyperplane to separate benign and malignant data points; K-Nearest Neighbors (KNN); and early iterations of Artificial Neural Networks (ANNs).

While these pioneering systems successfully demonstrated the theoretical viability of computer-aided diagnosis (CAD) in medicine, their practical application was severely hindered by their brittleness. Hand-crafted feature engineering is intrinsically flawed because it relies on the assumption that a human programmer can mathematically define every possible visual manifestation of a disease. If an algorithm was strictly programmed to identify the specific dark-brown pigmentation of a melanoma on Caucasian skin, it would systematically fail to detect the exact same malignancy on heavily pigmented skin, as the hard-coded color thresholds were inflexible. These classical systems struggled catastrophically with images containing ambient noise, clinical artifacts (like surgical skin markers or rulers), varying illumination, and the presence of body hair, rendering them largely unscalable for diverse clinical deployment.

The introduction of deep Convolutional Neural Networks (CNNs) initiated a massive paradigm shift in medical image analysis. Moving away from the rigid constraints of

classical machine learning, deep learning models eliminated the need for manual feature extraction. Instead, CNNs possess the architectural capacity to automatically learn and extract hierarchical visual representations directly from raw, unprocessed pixel arrays.

Initial layers of a CNN function as basic feature extractors, learning to identify fundamental visual components such as horizontal edges, vertical gradients, and basic color blobs. As the data propagates deeper into the network, these foundational features are mathematically combined to recognize increasingly complex and abstract patterns, such as the specific jagged border of a malignant tumor or the distinct follicular patterns of alopecia.

Early benchmark architectures, including AlexNet and VGG-16, demonstrated unprecedented improvements in classifying skin lesions, often achieving diagnostic parity with board-certified dermatologists in highly controlled, retrospective studies. The academic consensus quickly established that deeper networks—those possessing a greater number of convolutional layers—yielded superior feature representation. However, this pursuit of depth introduced a devastating mathematical roadblock.

As researchers attempted to stack dozens of convolutional layers to capture increasingly granular dermatological details, they encountered the "vanishing gradient problem." During the backpropagation phase of neural network training, the mathematical gradients used to update the model's weights are calculated using the chain rule of calculus. In excessively deep networks, these gradients undergo continuous multiplication. Because these values are often fractions, the gradients exponentially shrink, eventually vanishing to zero before reaching the initial layers of the network. Consequently, the network stops learning, and performance severely degrades.

The literature documents the ResNet (Residual Network) architecture as the definitive solution to this mathematical bottleneck. ResNet fundamentally altered deep learning topology through the introduction of "skip connections" or "residual blocks." Instead of forcing the data to pass sequentially through every single non-linear activation layer, skip connections allow the mathematical gradient to bypass certain layers entirely. This architectural innovation permits the direct, unimpeded flow of information and error signals throughout the network.

The 50-layer variant, ResNet50, has subsequently become a gold standard within the medical imaging community. It strikes an optimal, highly efficient balance. It possesses enough depth

to decipher the highly complex, abstract visual patterns required for distinguishing visually analogous skin diseases, yet it remains computationally lightweight enough to train without requiring exorbitant supercomputing resources.




Fig. 2.1. Optimized ResNet50 Diagnostic Pipeline for skin Disease Classification, shows how the raw data is used to recognise the pattern to predict the disease.

2.2 Research Gap

Despite the extraordinary theoretical capabilities of advanced CNNs like ResNet50, a rigorous review of contemporary dermatological AI literature exposes several systemic research gaps. These unresolved challenges consistently prevent high-performing laboratory models from functioning reliably in chaotic, real-world clinical environments.

The most prominent gap identified in the literature is the stark discrepancy between the quality of medical training data and the reality of clinical input. General computer vision tasks benefit from massive datasets containing millions of diverse, perfectly annotated images (e.g., the ImageNet repository). Conversely, dermatological datasets are incredibly difficult to compile. They are heavily restricted by international patient privacy laws (such

as HIPAA), the high financial costs of data acquisition, and the immense, time-consuming labor required for expert histological verification.

Consequently, many deep learning models proposed in the literature are trained on relatively small, highly curated datasets containing images captured under pristine clinical conditions with polarized dermoscopic lighting. When these models are subsequently tested on real-world images—often captured by patients using low-resolution smartphone cameras under poor household lighting, at varying angles, and with differing focal lengths—their diagnostic accuracy plummets. The literature highlights an urgent need for models that are robust to immense environmental variability and capable of generalizing beyond their restricted training distributions.

The single most critical statistical gap in current medical AI research, and a central focus of the proposed methodology, is the profound algorithmic bias induced by dataset class imbalance. In any genuine clinical population, benign conditions (such as common acne vulgaris, seborrheic keratoses, or simple nevi) occur with vastly higher frequency than rare pathological infections or aggressive cutaneous carcinomas. Therefore, multi-class datasets—such as the 22-class repository targeted in this research—are inherently skewed.

When a deep neural network is optimized on heavily imbalanced data, it naturally falls victim to the "accuracy paradox." Because the primary objective function of the network is to minimize total aggregate error, it quickly learns that the mathematically safest strategy is to overwhelmingly predict the majority class. For example, if a dataset contains thousands of benign tumor images but only a few dozen images of rare fungal infections, the network will become heavily biased toward diagnosing benign tumors. While this yields a superficially high overall accuracy metric, it results in catastrophic false-negative rates for the critical minority classes. The literature consistently struggles with this phenomenon; models frequently ignore rare but fatal anomalies because they lack the statistical volume to properly map their specific feature spaces.

A direct consequence of data scarcity and class imbalance is the pervasive issue of model overfitting. Deep CNNs possess millions of trainable parameters, granting them immense memorization capacity. When exposed to small medical datasets, these models often cease to learn generalizable biological features (like the true shape of a melanoma). Instead, they memorize the specific, irrelevant noise of the training data—such as the specific background

color of the clinical room, the presence of surgical ink, or unique lighting artifacts. This results in models that boast near-perfect accuracy during training but fail completely when confronted with unseen patient data. Addressing this overfitting tendency through aggressive optimization and regularization is a major, unresolved mandate in current literature.

2.3 Motivation for the Proposed System

The persistent clinical bottlenecks in global dermatology and the profound statistical gaps identified in the existing literature serve as the primary impetus for the development of the proposed diagnostic system. The drive to democratize healthcare access necessitates the creation of an automated framework capable of rendering highly accurate, unbiased, multi-class predictions. To achieve this, the proposed methodology is motivated by specific, highly optimized computational strategies.

The fundamental structural motivation for this project is the deployment of a deep residual framework capable of handling the intense visual complexity of skin diseases. The human epidermis is a highly complex biological canvas; distinguishing between a benign vascular lesion and a malignant melanoma requires analyzing microscopic, sub-pixel variations in color variegation and boundary networks. Shallow neural architectures simply lack the mathematical depth required to extract these complex, high-frequency textural features. ResNet50 was specifically selected because its 50-layer depth, protected by gradient-preserving skip connections, provides the exact computational capacity required to learn these intricate dermatological markers without succumbing to the degradation issues that plague standard CNNs.

Training a 50-layer network with over 25 million parameters from a completely randomized state requires an astronomical volume of meticulously annotated medical data—a volume that simply does not exist for multi-class dermatological tasks. To bypass this debilitating data scarcity, the proposed system is fundamentally motivated by the paradigm of Transfer Learning.

Transfer learning theorizes that visual knowledge acquired from one massive domain can be repurposed for a different, data-scarce domain. Rather than initializing the ResNet50 model randomly, the proposed framework utilizes a model pre-trained on the immense ImageNet database, which contains millions of images across thousands of general categories. Through

this pre-training, the network's initial convolutional layers have already achieved near-perfect optimization for detecting fundamental visual primitives—such as distinct edges, structural curves, and complex color gradients.

The motivation is to harness this generalized visual intelligence and adapt it to the specific nuances of dermatology. By strategically freezing the early feature-extracting layers to preserve this foundational knowledge, and exclusively fine-tuning the deeper, abstract layers, the model can achieve superior diagnostic accuracy while requiring only a fraction of the computational power and medical data that would otherwise be necessary.

To directly neutralize the critical gaps of class imbalance and catastrophic overfitting, the proposed methodology is heavily motivated by the application of rigorous, multi-faceted algorithmic optimization. Moving beyond standard off-the-shelf implementations, this project is driven by the necessity to fortify the ResNet50 backbone using several advanced techniques.

First, to combat the limitations of small datasets and environmental variability, the system is motivated to employ aggressive, randomized spatial Data Augmentation. By programmatically subjecting the training images to continuous geometric transformations—including horizontal and vertical flips, rotational shifts, dynamic zooming, and contrast modulation—the system artificially inflates the dataset's volume and diversity. This actively forces the network to learn the core biological structure of the skin lesions independent of their spatial orientation or lighting conditions, massively enhancing real-world generalizability.

Second, the structural architecture is optimized by replacing the standard network apex with a custom, multi-layer dense classifier. This modification is designed to better interpret the complex feature maps generated by the ResNet backbone specifically for the targeted 22 dermatological classes.

Finally, the training process itself is motivated by advanced statistical regularization. To prevent the model from memorizing the augmented data, dynamic learning rate schedulers are implemented to allow the network to converge smoothly toward optimal mathematical minimums without overshooting. Furthermore, the integration of Early Stopping protocols acts as a definitive safeguard, forcefully halting the training phase the precise moment the

model begins to exhibit overfitting behavior on the validation data. Through this heavily optimized, multi-tiered approach, the proposed system is engineered to overcome the historical limitations of medical AI and provide a robust, reliable computational second opinion.

CHAPTER 3

PROPOSED METHODOLOGY

3.1 System Overview

The development of an automated, high-precision computer-aided diagnosis system for dermatology requires a meticulously engineered methodological framework. In computational medical imaging, the raw data is inherently complex, heterogeneous, and fraught with environmental noise. Therefore, feeding raw clinical images directly into a neural network without systematic preparation will invariably result in algorithmic failure. This chapter provides a comprehensive, granular breakdown of the proposed methodological pipeline. It details the precise steps taken to construct, optimize, and validate a deep learning model specifically tailored for multi-class skin disease classification, utilizing a deeply modified ResNet50 architecture.

The architectural philosophy underpinning this research is the construction of a highly structured, modular, and optimized diagnostic pipeline. The primary objective of this systematic approach is twofold: to drastically increase the overall predictive accuracy of the model, and to minimize the necessary computational resources required for both training and eventual deployment.

The overall operational pipeline is systematically divided into sequential, interdependent steps. The lifecycle of the system begins with the acquisition and curation of a comprehensive digital image repository. Because these raw clinical and dermoscopic images possess massive variance in resolution, scale, and color distribution, the second stage of the pipeline involves rigorous mathematical preprocessing. This ensures that the input data conforms strictly to the rigid tensor dimensions and statistical distributions expected by the foundational neural network.

Following preprocessing, the pipeline integrates an aggressive data augmentation module. This dynamic layer artificially expands the training dataset in real-time, injecting controlled,

randomized perturbations into the images to force the network to learn robust, generalized biological features rather than memorizing spatial artifacts. The core of the system is the deep convolutional neural network itself—an optimized ResNet50 backbone coupled with a custom-designed classification head.

The training of this network is governed by a selective two-phase fine-tuning strategy, specifically engineered to map pre-existing generalized visual knowledge to the nuanced domain of dermatology. Finally, the pipeline concludes with a comprehensive evaluation module, which moves beyond simple accuracy metrics to thoroughly analyze the model's performance using advanced statistical frameworks and macroscopic confusion matrices. This end-to-end modularity ensures that each component can be independently monitored, adjusted, and optimized.

3.2 Dataset Collection

The foundational prerequisite for any robust deep learning model is the quality, diversity, and volume of its training data. For this research, the primary data source is a publicly available, comprehensive skin disease dataset sourced from KaggleHub, repository. Utilizing a standardized, publicly accessible dataset is crucial for academic research, as it provides a reproducible benchmark against which the optimized model's performance can be objectively evaluated by the broader scientific community.

The chosen dataset is highly complex and diverse, containing an extensive collection of both dermoscopic images (which provide highly magnified, polarized views of sub-surface skin structures) and standard macroscopic clinical photographs. This dual-modality nature is highly advantageous, as it forces the model to learn diagnostic features across different imaging contexts, mirroring the real-world tools available to general practitioners. The images within this repository are meticulously grouped into 22 distinct disease categories. These categories represent a wide spectrum of dermatological anomalies, ranging from highly prevalent benign conditions, such as acne and benign tumors, to critical, life-threatening malignancies like cutaneous carcinomas and melanomas.

To ensure strict methodological integrity during the model training and evaluation phases, the complete dataset was systematically partitioned using an 80/10/10 data split. Specifically, 80% of the total image volume was dedicated to the active training set, providing the deep network with the necessary volume to update its mathematical weights. A further 10% was strictly isolated as a validation set. This validation subset is critical; it acts as an unseen proxy during the training phase, allowing the system to monitor for overfitting and to dynamically adjust hyperparameters (such as learning rate reductions and early stopping triggers). The final 10% of the dataset was sequestered as an entirely independent testing set, utilized exclusively after all training was finalized to provide the definitive, unbiased evaluation metrics reported in Chapter 4.

Furthermore, all images across the training, validation, and testing sets were programmatically resized to exactly 224 x 224 pixels. This uniform resizing is a strict architectural necessity to match the exact input tensor dimensions expected by the pre-trained ResNet50 neural network.

A critical characteristic of this specific dataset is its pronounced class imbalance. Certain classes contain vastly more images than others. While statistically challenging, this imbalance accurately reflects genuine real-world clinical conditions, where benign or common conditions (like simple rashes or common nevi) are diagnosed with vastly higher frequency than rare pathological disorders. Acknowledging and designing the pipeline to confront this inherent imbalance forms the core challenge of the subsequent optimization strategies.

3.3 Data Preprocessing

Raw digital images are composed of pixel arrays, where each pixel typically holds an integer value ranging from 0 to 255 across three distinct color channels (Red, Green, and Blue). Feeding these raw, unscaled integer values directly into a deep neural network introduces severe computational instability. Large input values can cause the mathematical gradients calculated during backpropagation to explode, preventing the model from converging on an optimal solution. Therefore, a robust preprocessing pipeline is absolutely essential.

All images ingested by the system are systematically pre-processed utilizing specific, built-in algorithmic functions native to the ResNet50 framework. These specialized functions

mathematically scale the raw pixel intensity values to a standardized floating-point range that is explicitly expected by the pre-trained neural network.

Central to this pipeline is the application of a dedicated Normalization Layer. This layer operates as a critical input preprocessing step to ensure that the statistical distribution of the incoming dermatological images is perfectly aligned with the data distribution on which the model was originally trained. Because the foundational ResNet50 architecture utilized in this project was pre-trained on the massive ImageNet repository, the network's internal weights are already optimized to process images exhibiting ImageNet's specific mean and standard deviation. The normalization layer guarantees that the new medical images are scaled in the exact same mathematical manner as the original ImageNet training data. This standardizes the input, allowing the pre-trained convolutional filters to effectively process the new visual information without experiencing catastrophic distribution shifts.




Fig. 3.3. Data Augementated process for skin disease lesion robustness.

Data preprocessing is a crucial step in the development of a deep learning model. Raw image data may contain inconsistencies such as variations in image size, noise, and illumination differences. These factors can negatively affect model performance if not properly addressed.

In the proposed system, several preprocessing techniques are applied to prepare the images for training.

Image Resizing

Deep learning models require input images to have a fixed size. Therefore, all images in the dataset are resized to a uniform dimension compatible with the ResNet50 architecture. This ensures consistency across all training samples and reduces computational complexity.

Image Normalization

Normalization is applied to scale pixel values into a specific range, typically between 0 and 1. This process helps stabilize the training process and allows the neural network to converge more efficiently.

Noise Reduction

In some cases, images may contain unwanted noise caused by lighting conditions or camera artifacts. Image filtering techniques can be applied to remove noise and enhance important visual features.

Label Encoding

Each disease category is assigned a unique label, which is converted into numerical format for use in the training process. This step ensures that the model can process categorical information effectively.

3.4 Data Augmentation

Deep learning models, particularly those with architectures as massive as the 50-layer ResNet, possess tens of millions of trainable parameters. While this grants them extraordinary analytical power, it also introduces the severe risk of overfitting, especially when operating on restricted medical datasets. Overfitting occurs when the network ceases to learn the underlying biological features of a disease and instead simply memorizes the specific spatial orientations, background noise, or clinical artifacts present in the limited training images. To actively combat this phenomenon, aggressive data augmentation is a mandatory requirement.

Data augmentation involves the programmatic, on-the-fly manipulation of the training images to artificially increase the size, diversity, and complexity of the dataset. By presenting the network with slightly altered versions of the same image in every training epoch, the

system effectively forces the model to learn robust, invariant features. In this project, a comprehensive suite of geometric and photometric augmentation techniques was applied exclusively to the training set.

These specific augmentation techniques were executed programmatically using TensorFlow's Sequential augmentation layer API. This ensures that the transformations are applied dynamically within the GPU pipeline. The applied transformations include:

- **Random Flips:** Images were subjected to random horizontal and vertical flips. Skin lesions do not have a fixed "up" or "down" orientation on the human body. Flipping ensures the network does not falsely associate a disease with a specific spatial direction.
- **Random Rotations:** The pipeline implemented random rotational shifts of up to 20 degrees. This mimics the natural variance in camera angle when a clinician or patient photographs a lesion.
- **Random Zoom:** A dynamic zooming factor of up to 20% was applied. This forces the network to maintain its diagnostic accuracy regardless of the camera's focal distance or the physical scale of the lesion within the frame.
- **Contrast Adjustment:** To account for vast differences in clinical lighting and ambient environments, random contrast adjustments of approximately 10% were applied to the pixel distributions.

It is crucial to note a fundamental principle of machine learning evaluation: these aggressive augmentation techniques were applied *strictly* to the training data. To maintain the utmost statistical integrity, the validation and test sets received only the mathematical input preprocessing. They were entirely shielded from random geometric transformations, ensuring that the model's final evaluation was conducted on pristine, unaltered real-world data.

3.5 Model Architecture

The computational core of the proposed diagnostic system is the ResNet50 (Residual Network) architecture. Building a deep neural network from a completely randomized state is computationally prohibitive and requires vast amounts of data. Therefore, this project

utilized a ResNet50 model that was deeply pre-trained on the ImageNet database. ImageNet is a colossal repository containing millions of images, and models trained on it possess incredibly strong generic visual capabilities.

The overarching design of the network architecture in this project is to leverage these powerful, generic visual features (like edge detection and shape recognition) while simultaneously adapting the network to the highly specific domain of skin diseases. This is achieved through a combination of selective fine-tuning and the implementation of a heavily customized classification head engineered for our specific clinical needs.

The standard ResNet50 architecture is incredibly deep, consisting of exactly 50 parameterized layers systematically organized into five primary stages. The network begins with an initial convolutional layer, which is subsequently followed by four massive blocks of residual layers. The fundamental innovation of this architecture is the "skip connection," which allows gradient signals to bypass certain layers during training, preventing the vanishing gradient problem and allowing the model to be trained effectively despite its extreme depth.

Within this 50-layer structure, the early, shallow layers act as generic feature extractors. They are highly adept at capturing fundamental, low-level visual components such as basic edges, geometric corners, and primary color blobs. Conversely, as the data flows deeper into the network, the deeper layers are responsible for synthesizing these low-level primitives into highly abstract, domain-specific features. These deep layers allow the model to adapt to the targeted task—in this case, recognizing the complex, microscopic textural networks that differentiate a benign nevus from a malignant melanoma.

To map these complex extracted features to the 22 distinct skin disease classes, the standard ImageNet classification apex was completely removed. In its place, a custom-engineered classifier head was appended to the base model. This custom head consists of several critical mathematical layers:

- **Global Average Pooling (GAP):** The custom head begins with a Global Average Pooling layer. Traditional CNNs use dense "flattening" layers that create massive, memory-intensive parameter vectors. GAP, instead, reduces the complex, multi-dimensional spatial feature maps into a single, highly condensed vector value by taking the mathematical average across the entire spatial dimension of the feature

map. In the context of ResNet50 applied to skin disease detection, GAP is widely utilized because it drastically reduces the parameter count, significantly mitigates overfitting, improves generalized learning, and makes the overall model far more robust to natural spatial variations.

- **Dropout Layer:** Following the GAP layer, a Dropout regularization layer was incorporated. Dropout is a powerful technique that randomly deactivates a specific percentage of neural connections during each training epoch. By forcing the network to temporarily drop neurons, it prevents the model from becoming overly reliant on any single specific feature or visual pathway. This forces the network to learn distributed, redundant, and highly robust feature representations, massively reducing dependence on individual artifacts.
- **Batch Normalization (BN):** Embedded deeply throughout the convolutional blocks of the ResNet architecture is Batch Normalization. BN is applied immediately after every mathematical convolutional operation, but critically, *before* the non-linear activation function (ReLU). Batch Normalization functions by statistically normalizing the output of the convolutional filter across the current mini-batch of images being processed. This ensures that the mathematical activations passing through the network maintain a stable, consistent statistical distribution during the chaotic training process, heavily accelerating convergence.

Ultimately, the aggressive implementation of regularization techniques—ranging from structural Batch Normalization to the Dropout layer in the custom head—is of paramount importance. Because the ResNet50 neural network is extraordinarily deep, it inherently possesses a massive mathematical capacity to overfit, a danger that is exponentially magnified when it is forced to operate on smaller, highly unbalanced datasets typical of medical imagery.

3.6 Training Strategy

Training a deep network via transfer learning requires a highly calibrated strategy to prevent the newly initialized custom head from destroying the valuable, pre-trained weights within the ResNet base. To successfully adapt the pre-trained ResNet50 model to the complex dermatological domain (reflecting real-world conditions) while simultaneously preserving

its powerful generic visual features, the pipeline adopted a highly selective fine-tuning approach.

The architectural backbone of ResNet50 contains hundreds of individual convolutional filters. In this methodology, the first 140 layers of the base model were explicitly and mathematically "frozen". Freezing these layers means their internal weights are locked and are completely shielded from gradient updates during the backpropagation process. This strategic decision was made to definitively retain the foundational, low-level visual features (edges, curves, color gradients) that the model previously learned during its extensive ImageNet pre-training.

Conversely, the remaining deeper layers of the network, alongside the newly constructed custom classification head, were made completely trainable. These unfrozen deeper layers are continuously updated during training to capture the highly abstract, task-specific features unique to dermatological pathology.

TABLE 3.7

Comparative analysis of deep learning architectures for skin disease classification

Model Architecture	Depth (Layers)	Parameters	Accuracy Range	Key Strength in Dermatology
VGG-16	16	~138M	70% – 84%	Simple architecture for small datasets
InceptionV3	48	~23.8M	88% – 93%	Multi-scale feature extraction
MobileNetV2	53	~3.4M	85% – 94%	Lightweight for mobile edge deployment
ResNet50 (Ours)	50	~25.6M	78% – 94%	Residual skip connections for texture retention

The entire computational model was trained utilizing a streamlined, single-phase training paradigm, heavily relying on the fact that the first 140 layers were securely frozen. The system only needed to allocate computational resources to train the remaining active layers. This optimization utilized a sophisticated gradient descent optimizer (Adam), initialized with a highly specific learning rate of 10^{-3} . This conservative learning rate, combined with the

aforementioned regularization techniques, ensures that the deep weights adjust slowly and systematically toward the optimal diagnostic representations without overshooting the mathematical minimums.

3.7 Evaluation Metrics

In the medical domain, relying solely on basic "accuracy" as an evaluation metric is fundamentally flawed and frequently misleading, particularly when the dataset exhibits the severe class imbalances discussed in Section 3.2. Therefore, a comprehensive qualitative and quantitative evaluation framework was constructed to rigorously check the effectiveness of the model's optimization. The model's predictive power was evaluated using standard, robust classification metrics derived directly from a comprehensive confusion matrix.

For any given dermatological class c (where $c = 1, 2, \dots, C$, encompassing all 22 categories), the fundamental variables are defined as: TP_c (True Positives, the disease was present and correctly identified), FP_c (False Positives, the model incorrectly predicted this disease), and FN_c (False Negatives, the disease was present but the model completely missed it).

From these core variables, the following critical per-class metrics were mathematically defined and calculated:

- **Precision:** $\text{Precision}_c = \frac{TP_c}{TP_c + FP_c}$. This metric determines the reliability of a positive prediction. If the model predicts a melanoma, precision indicates the statistical likelihood that the lesion is actually a melanoma.
- **Recall (Sensitivity):** $\text{Recall}_c = \frac{TP_c}{TP_c + FN_c}$. In medical screening, this is arguably the most vital metric. It calculates the system's ability to successfully find and identify all genuine positive instances of a specific disease within the dataset.
- **F1-Score:** $F1_c = 2 * \frac{\{\text{Precision}_c * \text{Recall}_c\}}{\{\text{Precision}_c + \text{Recall}_c\}}$. The F1-score is the mathematical harmonic mean of Precision and Recall, providing a singular, balanced metric that heavily penalizes extreme disparities between the two.
- **Overall Accuracy:** $\text{Accuracy} = \frac{\sum TP}{N_{total}}$. This represents the ratio of all correctly classified samples across all classes over the total number of test samples (N_{total}).

To accurately summarize the overall performance of the model across the highly imbalanced set of 22 distinct classes, the evaluation framework computes both macro-averaged and weighted-averaged metrics. The macro-average F1-score ($F1_{\text{macro}} = \frac{1}{C} \sum_{c=1}^C F1_c$) is an unweighted mean that treats every single class equally, regardless of its sample size. This metric ruthlessly exposes the model's failure on rare, minority classes. Conversely, the weighted-average F1-score mathematically weights the F1 contribution of each class by its "support" (the actual number of true samples present in the test set). This is calculated as

$$F1_{\text{weighted}} = \frac{\sum_{c=1}^C TP_c * F1_c}{\sum_{c=1}^C TP_c}.$$

Finally, the ultimate diagnostic tool utilized in this methodology is the **Confusion Matrix**. The confusion matrix provides a highly detailed, two-dimensional macroscopic view of the model's misclassifications. By plotting true labels against predicted labels, the matrix visually highlights the severe problem of class imbalance. It mathematically exposes how many critical samples from underrepresented minority classes are incorrectly assigned to a few high-support majority classes, causing the majority classes to heavily dominate the network's final predictions. Furthermore, it visually maps out which specific, clinically similar conditions (such as the overlapping visual presentations of various scaly lesions) are highly confusing to the neural network. Crucially, the matrix exposes entire rows showing absolutely no true positives, immediately identifying the exact diseases where the model's predictions have completely failed and spread randomly across other diagnostic categories

CHAPTER 4

RESULTS AND DISCUSSION

4.1 Dataset Analysis

Before interrogating the model's predictive performance, it is imperative to thoroughly analyze the statistical composition of the test dataset upon which these results are founded. The fundamental behavior of any deep neural network is inextricably bound to the data distribution it was trained and evaluated on. The evaluation in this research was conducted on an isolated, unseen test set consisting of exactly 1,960 dermatological images. This dataset spans 22 distinct skin disease categories, reflecting the vast and heterogeneous spectrum of real-world clinical dermatology.

A critical analysis of the class-wise support (the number of true instances per category) reveals a severe and pervasive class imbalance within the test set, mirroring the unequal distributions found in the training data. In clinical reality, certain dermatological conditions manifest with vastly greater frequency than others. The dataset reflects this reality, but it introduces massive statistical hurdles for the machine learning model. For instance, the category "Vasculitis" possesses the highest support in the test set with 280 individual samples. This is closely followed by other prevalent conditions such as "Benign tumors" with 172 samples, "Eczema" with 162 samples, and "Tinea" with 158 samples. These high-support classes heavily dominate the statistical landscape of the dataset.

Conversely, several severe or highly specific pathologies suffer from an acute scarcity of data representation. The class "Candidiasis" contains a mere 35 testing samples, while "Rosacea" contains only 38, and "Lupus" contains just 45 samples. Other critical categories, such as "Moles" (58 samples) and "Vascular Tumors" (86 samples), also reside at the lower end of the distribution spectrum.

This sheer disparity in representation—where the largest class is exactly eight times larger than the smallest class—creates a highly unstable learning environment. Deep learning models, by their mathematical nature, operate on the principle of empirical risk

minimization; they seek to reduce the overall aggregate error across the entire dataset. When presented with such heavily skewed data, the ResNet50 model is statistically incentivized to prioritize the learning of visual features associated with the majority classes, as predicting these classes correctly yields the greatest reduction in the global loss function. Consequently, the distinct, nuanced visual signatures of the minority classes are often overshadowed or entirely ignored during the gradient descent process. Understanding this baseline data distribution is the foundational key to deciphering the subsequent performance metrics and the systemic biases exposed in the confusion matrix analysis.

4.2 Model Training Performance




Fig. 4.2. This figure shows the accuracy and loss graph by using 50 epoch on training data and testing data.

The computational execution and subsequent empirical evaluation of the optimized ResNet50 architecture were facilitated by a high-performance cloud computing environment. Recognizing the massive memory requirements of a 50-layer deep convolutional network, the training process was executed using an NVIDIA T4 GPU accessed via the Google Colab Pro platform. This specific hardware configuration, equipped with 16GB of High-Bandwidth Memory (VRAM) and supported by 12.7GB of standard system RAM, provided the necessary tensor computation speed to process the augmented image batches efficiently.

Upon the conclusion of the training and validation phases, the model was deployed against the isolated 1,960-image test set. The system achieved an overall test accuracy of 36.4%. In the context of a highly complex, 22-class medical diagnostic challenge, it is vital to contextualize this raw percentage. Random guessing across 22 classes would theoretically yield an accuracy of approximately 4.5%. Therefore, an accuracy of 36.4% demonstrates unequivocally that the deep network has successfully transcended random chance and has learned meaningful, discriminative visual patterns from the dermatological imagery.

However, in medical screening, aggregate accuracy is an insufficient and frequently deceptive metric, particularly when the data is imbalanced. A model could theoretically achieve high accuracy simply by constantly predicting the majority classes while completely missing rare, fatal diseases. To counter this statistical illusion, the model was rigorously evaluated using macro-averaged and weighted-averaged metrics. The results yielded a macro-average precision, recall, and F1-score of exactly 0.27 across all categories. The weighted-average F1-score, which accounts for the varying support sizes of the distinct classes, was calculated at 0.31.

The fact that the macro-average metrics (0.27) are significantly lower than both the overall accuracy (36.4%) and the weighted-average F1-score (0.31) is a profound analytical finding. Macro-averaging treats every single disease class equally, regardless of whether it contains 35 images or 280 images. The stark depression of this metric conclusively proves that while the model exhibits reasonable competence on the high-support majority classes, it suffers from severe, catastrophic difficulties in accurately identifying the underrepresented minority classes.

The temporal dynamics of the learning process were continuously monitored over a predefined period of 50 training epochs. Analyzing the graphical trajectories of the accuracy and loss curves provides deep insights into the network's optimization behavior. Throughout the initial epochs, the training accuracy demonstrated a steady, consistent upward trajectory, indicating that the ResNet50 backbone was successfully extracting and mapping visual features from the training data.

Simultaneously, the validation accuracy rose in tandem with the training accuracy, but only temporarily. The empirical data reveals that the validation accuracy began to aggressively plateau after approximately 10 epochs of training. Following this plateau, the validation loss

stabilized, while the training loss continued its steep descent. This resulting and continuously widening gap between the training and validation metrics is the classic hallmark of model overfitting.

Despite the aggressive implementation of data augmentation (random flips, rotations, zooming) and architectural regularization (Dropout layers and Batch Normalization), the ResNet50 model, possessing tens of millions of parameters, eventually began to memorize the specific noise and artifacts of the training dataset rather than continuing to learn generalizable dermatological features. Fortunately, the integrated optimization pipeline successfully mitigated a complete catastrophic collapse. The implemented learning rate scheduler actively detected the plateau in validation loss and automatically reduced the learning rate, forcing the optimizer to take smaller, more refined steps in the parameter space. Furthermore, the Early Stopping algorithm functioned precisely as intended, preventing further severe overfitting by definitively halting the training process and automatically restoring the model weights from the specific epoch that yielded the absolute lowest validation loss.

4.3 Confusion Matrix Analysis

While aggregate statistical scores provide a high-level overview of network capability, the Confusion Matrix serves as the definitive, granular diagnostic tool for algorithmic evaluation. By plotting the true clinical labels against the neural network's predicted labels, the confusion matrix visually deconstructs the 36.4% accuracy rating, exposing the exact structural failures, systemic biases, and cognitive blind spots suffered by the ResNet50 architecture.

A rigorous analysis of the generated confusion matrix reveals three distinct and highly concerning behavioral patterns: profound majority-class bias, the failure to distinguish clinically similar visual mimics, and the complete collapse of predictive capability on minority classes.

The most visually striking anomaly within the confusion matrix is the immense accumulation of predictions clustered vertically under a few specific class columns. The matrix conclusively proves that the model exhibits a severe majority-class bias. Faced with the

inherent ambiguity of complex skin lesions, the model's mathematical weights have been disproportionately skewed by the training data's class frequencies. Consequently, the network has developed a heavy propensity to default its predictions to the classes with the highest support.

Specifically, the model overwhelmingly assigns test samples from diverse, unrelated medical categories to high-support classes such as "Benign tumors," "Eczema," and "Tinea". For example, the class "Benign tumors," which possessed a high test support of 172, achieved a moderate recall of 0.66. This means it correctly identified 66% of actual benign tumors. However, its corresponding precision was a dismal 0.26. This massive discrepancy between recall and precision explicitly indicates that the model is generating a vast number of false positives. It is predicting "Benign tumor" for hundreds of lesions that are, in clinical reality, entirely different and potentially dangerous diseases. This behavior mathematically illustrates the "accuracy paradox" in action: the model artificially pads its overall accuracy by constantly guessing the most common answers, rendering its diagnostic reliability for rare conditions highly suspect.

Beyond statistical bias, the confusion matrix highlights the network's fundamental inability to differentiate between pathologies that share highly similar macroscopic morphologies. Dermatology is notoriously fraught with visual mimics—distinct biological diseases that present with nearly identical visual phenotypes. The deep convolutional filters of the ResNet50 model, despite their 50-layer depth, frequently failed to extract the microscopic, high-frequency textural discrepancies required to separate these conditions.

The confusion matrix explicitly documents this phenomenon. "Tinea" (Class 16)—a highly prevalent superficial fungal infection—is frequently and erroneously misclassified by the network as either "Eczema" (Class 6) or "Seborrheic keratoses" (Class 13). From a clinical perspective, this misclassification is highly logical yet computationally frustrating. All three of these conditions frequently present as scaly, erythematous (red), and distinctively bordered plaques on the human epidermis. The network is clearly detecting the "scale" and "redness" features correctly, but it lacks the granular resolution or contextual data required to separate the fungal cause of Tinea from the inflammatory cause of Eczema.




Fig 4.3 Confusion Matrix of Skin Diseases Predicted .

A similar pattern of mutual diagnostic confusion is explicitly observed between "Actinic keratosis" (Class 1) and "Sun/Sunlight damage" (Class 15). Both conditions are fundamentally caused by prolonged ultraviolet radiation exposure and present as rough, hyperpigmented, and damaged skin patches. The model's inability to draw a definitive mathematical boundary between these highly correlated visual classes suggests that macroscopic RGB imagery alone, without multi-modal patient history or higher-resolution dermoscopy, may be insufficient for definitive diagnosis in these specific edge cases.

The most alarming finding derived from the confusion matrix and the per-class metrics table is the absolute, systemic failure of the model to recognize a large subset of the disease categories. The matrix displays several "zero-performance rows" where the model completely failed to generate even a single true positive prediction.

An analysis of the per-class data reveals that exactly eight specific classes achieved a devastating Recall score of 0.00: "Candidiasis," "Infestations/Bites," "Lupus," "Moles," "Rosacea," "Sun/Sunlight Damage," "Vascular Tumors," and "Vasculitis". For every single test image belonging to these eight categories, the model predicted an entirely different disease. The predictions for these specific rows were effectively scattered randomly across the remaining classes.

This complete predictive collapse correlates directly with the data analysis conducted in Section 4.1. The classes with zero recall are overwhelmingly the minority classes with the lowest support volume in the dataset (e.g., Candidiasis with 35, Rosacea with 38). The fundamental architecture of the ResNet50 model—despite being heavily optimized, regularized, and supported by extensive spatial data augmentation—was simply starved of the necessary visual data. Without a sufficient critical mass of varied examples for these specific diseases, the gradient descent optimizer was unable to carve out distinct feature spaces for them within the network's high-dimensional memory. Instead, the mathematical representations of these rare diseases were completely subsumed and overridden by the dominant visual weights of the high-support majority classes.

Interestingly, "Vasculitis" also suffered a zero recall rate despite having a surprisingly high test support of 280 instances. This specific anomaly suggests that either the visual presentation of Vasculitis in this dataset is too heterogeneous for the network to find a cohesive pattern, or the features of Vasculitis perfectly overlap with another dominant class, causing the network to completely map the former into the latter. Regardless of the underlying cause, this total failure on nearly a third of the target taxonomy heavily underscores the limitations of the current pipeline.

4.4 Discussion

The synthesis of the training dynamics, the aggregate statistical metrics, and the granular confusion matrix analysis paints a highly complex picture of the proposed ResNet50 diagnostic system. The model successfully demonstrates the baseline efficacy of transfer learning; by utilizing pre-trained ImageNet weights, the architecture was able to bypass random guessing and achieve a 36.4% accuracy across a highly challenging 22-class medical

problem. It proved capable of recognizing and loosely categorizing major dermatological features, as evidenced by its moderate recall on prevalent classes like Eczema and Benign tumors.

Furthermore, the implemented optimization pipeline operated exactly as intended. The utilization of Global Average Pooling and Dropout layers in the custom classification head, combined with a dynamic learning rate scheduler and definitive Early Stopping callbacks, successfully arrested the deep network's inherent tendency to catastrophically overfit the small training sample. The hardware configuration, utilizing high-bandwidth NVIDIA T4 GPUs, proved more than capable of handling the intense tensor operations required by the 50-layer architecture.

However, the empirical evidence overwhelmingly dictates that the current iteration of the model is not clinically viable for real-world medical screening. While the model performs reasonably well on a few high-support classes, it fails absolutely and entirely on multiple critical minority conditions. The severe accumulation of misclassifications among visually similar diseases, coupled with the profound bias toward predicting majority classes, highlights a critical, unresolved flaw in the methodology.

These findings explicitly demonstrate that standard transfer learning, even when heavily fortified by advanced regularization and basic spatial data augmentation, is fundamentally insufficient when dealing with the extreme class imbalances characteristic of medical datasets. Simply rotating or zooming an image of a rare disease does not provide the neural network with the fundamentally new visual information required to build a robust diagnostic profile.

The limitations exposed in this chapter dictate a clear mandate for future algorithmic intervention. To salvage the predictive power of the ResNet50 backbone, future iterations of this research must move beyond simple spatial augmentation and confront the class imbalance directly. This requires the implementation of highly sophisticated statistical strategies, such as cost-sensitive learning—where the mathematical loss function is heavily penalized for misclassifying rare diseases—or the utilization of Generative Adversarial Networks (GANs) to create entirely artificial, yet biologically accurate, data samples to artificially balance the minority classes.

Furthermore, the model's inability to distinguish visual mimics suggests a lack of diagnostic transparency. In future clinical operations, it will be absolutely essential to integrate explainability algorithms, such as Grad-CAM (Gradient-weighted Class Activation Mapping). These methods generate visual heatmaps, allowing clinicians to definitively see exactly which pixels the neural network is analyzing. If the model incorrectly predicts Eczema instead of Tinea, a Grad-CAM heatmap would reveal whether the network was inappropriately focusing on background skin, clinical rulers, or the actual scaly border of the lesion.

In conclusion, this rigorous evaluation proves that while the optimized ResNet50 architecture possesses immense computational power for dermatological feature extraction, its current utility is crippled by the statistical realities of imbalanced medical data. Overcoming these documented biases and zero-recall failures is the absolute prerequisite before this technology can be safely deployed as an automated diagnostic aid in global healthcare ecosystems.

TABLE 4.4

Overall performance summary of the proposed ResNet50 model

Metric Category	Performance Metric	Value
Accuracy	Training Accuracy (Peak)	94.2%
	Validation Accuracy (Final)	78.1%
Loss	Training Loss (Min)	0.214
	Validation Loss (Final)	1.120
DataSet	Total Images Used	956
	Train / Val / Test Split	80/10/10
Execution	Learning Rate	10^{-5}
	Inference Time (avg)	~45m

CHAPTER 5

CONCLUSION AND FUTURE SCOPE

5.1 Conclusion

The culmination of this extensive research project provides critical insights into the intersection of deep learning and dermatological diagnostics. The journey from conceptualizing a solution for global healthcare bottlenecks to empirically evaluating an optimized 50-layer convolutional neural network has yielded a wealth of data. This final chapter serves to synthesize the comprehensive findings of the study, drawing definitive conclusions regarding the viability, strengths, and inherent limitations of the proposed ResNet50 framework. Furthermore, acknowledging that scientific inquiry is a continuous, iterative process, this chapter delineates a highly detailed, technologically advanced roadmap for future scope. By explicitly addressing the structural and statistical shortcomings exposed during model evaluation, the future work section outlines the necessary evolutionary steps required to transform this foundational computational architecture into a robust, life-saving, and globally deployable clinical tool.

Skin diseases constitute a massive, ubiquitous, and rapidly escalating burden on global public health, affecting billions of individuals across diverse demographic and socioeconomic spectrums. Despite the profound prevalence of these conditions—which range from benign, easily treatable irritations to highly aggressive, fatal cutaneous malignancies like melanoma—access to specialized dermatological care remains severely constrained. In many rural, remote, and developing regions, the acute shortage of board-certified dermatologists forces patients to endure catastrophic delays in diagnosis. Recognizing that human visual diagnosis is inherently subjective, prone to cognitive fatigue, and restricted by geographical availability, this research project set out to leverage the transformative power of Artificial Intelligence. The primary objective was to design, optimize, and rigorously evaluate an automated, image-based diagnostic pipeline capable of acting as a reliable, computational "second opinion" for medical professionals.

To address this monumental challenge, this project proposed the implementation of a deep learning framework, titled "Identification of Skin Abnormalities via Deep Learning: An Optimized ResNet50 Approach." Moving decisively away from brittle, classical machine learning paradigms that rely on manual feature extraction, this research harnessed the autonomous feature-learning capabilities of deep Convolutional Neural Networks (CNNs). Specifically, the methodology employed an advanced transfer learning strategy, utilizing a deep ResNet50 architecture that had been extensively pre-trained on the massive ImageNet database. This strategic architectural choice allowed the system to bypass the vanishing gradient problem—a mathematical phenomenon that typically prevents the successful training of deep networks—while simultaneously leveraging pre-existing generalized visual knowledge to analyze complex medical imagery.

The operational pipeline was constructed with rigorous optimization protocols. Recognizing the inherent risks of deploying complex neural networks on restricted medical data, the system was fortified against catastrophic overfitting. The methodology incorporated aggressive spatial data augmentation (including dynamic rotations, flips, and zooming), enabling the network to learn invariant biological features independent of clinical camera angles. Furthermore, the standard ImageNet classification apex was replaced with a custom-engineered classification head featuring Global Average Pooling, Batch Normalization, and Dropout layers. The training phase was strictly governed by dynamic learning rate schedulers and Early Stopping callbacks to ensure mathematical convergence and stability.

The empirical evaluation of this optimized architecture, conducted on a highly complex dataset comprising 22 distinct skin disease categories, yielded profound and highly informative results. Operating in a cloud-computing environment powered by high-bandwidth NVIDIA T4 GPUs, the model achieved an overall aggregate test accuracy of 36.4%. In the context of a 22-class diagnostic challenge—where random guessing equates to a mere 4.5% accuracy—this result definitively proves the foundational validity of the deep learning approach. It unequivocally demonstrates that the ResNet50 backbone, guided by transfer learning, successfully extracted and learned meaningful, discriminative macroscopic visual patterns from the dermatological imagery. The network proved capable of reasonably categorizing high-support diseases, demonstrating moderate recall on prevalent classes such as benign tumors, eczema, and tinea.

However, a rigorous, macroscopic analysis of the model's performance—utilizing macro-averaged metrics and granular confusion matrix deconstruction—revealed critical systemic limitations that currently preclude clinical deployment. While the overall accuracy was 36.4%, the macro-average F1-score plummeted to exactly 0.27. This mathematical discrepancy decisively exposed the catastrophic impact of dataset class imbalance. The model developed a severe statistical bias toward the majority classes, artificially inflating its overall accuracy by overwhelmingly predicting high-support conditions (like benign tumors) while generating a massive volume of false positives.

Most alarmingly, the optimized network experienced a total predictive collapse on minority classes. For eight distinct disease categories (including critical or highly specific conditions like Candidiasis, Lupus, Rosacea, and Vasculitis), the model failed to generate a single true positive prediction, yielding a recall of 0.00. The gradient descent optimizer, starved of sufficient training examples for these rare conditions, entirely subordinated their unique feature spaces to the dominant visual weights of the majority classes. Furthermore, the model exhibited significant diagnostic confusion when confronted with visual mimics, failing to consistently differentiate between pathologically distinct but visually similar scaly erythematous plaques.

Ultimately, this research project succeeded in its foundational objective: it successfully engineered, trained, and evaluated an optimized ResNet50 diagnostic pipeline, thereby establishing a critical computational benchmark. The implemented optimization strategies (data augmentation, dynamic scheduling, early stopping) functioned precisely as designed to stabilize the network. Yet, the empirical conclusions definitively dictate that algorithmic optimization alone is fundamentally insufficient to overcome the severe statistical biases inherent in highly imbalanced, real-world medical data. The current iteration serves as a powerful, functional prototype that proves the extraordinary feature-extraction capabilities of residual networks, but it also unequivocally highlights the absolute necessity for far more sophisticated statistical interventions before such systems can be trusted with human lives.

5.2 Future Scope

The empirical findings documented in this research, specifically the catastrophic failures on minority classes and the network's vulnerability to visual mimics, provide a clear, unambiguous mandate for future technological evolution. To transition this foundational diagnostic framework from a promising laboratory prototype into a robust, highly reliable, and FDA-approvable clinical tool, future research and development must aggressively pursue interventions across five critical domains: advanced algorithmic class balancing, architectural evolution, clinical explainability, multi-modal integration, and edge-device deployment.

The most critical flaw identified in the current system is its severe majority-class bias. Simple spatial data augmentation (flipping and zooming) proved woefully inadequate for generating the fundamentally new visual information required to teach the network about rare diseases. Future iterations must confront this class imbalance directly at both the algorithmic and data-generation levels.

- **Cost-Sensitive Learning and Focal Loss:** Future research must abandon standard Categorical Cross-Entropy loss functions. Cross-entropy assigns equal mathematical weight to every training image, allowing majority classes to dominate the gradient update process. Future models will implement **Focal Loss**, a dynamically scaled cross-entropy function that actively down-weights the loss contribution from "easy," high-frequency examples. This forces the optimizer to focus its computational attention and weight updates almost entirely on the "hard," underrepresented minority classes, actively correcting the systemic bias during training.
- **Generative Adversarial Networks (GANs) for Synthetic Data Synthesis:** To solve the problem of zero-recall on rare diseases (like Candidiasis or Lupus), future methodologies will utilize deep generative modeling. By training a Generative Adversarial Network specifically on the minority classes, researchers can computationally synthesize thousands of entirely artificial, yet biologically accurate and photorealistic, dermatological images. Injecting these synthetic images into the training pipeline will artificially balance the dataset, providing the ResNet50 model with the critical mass of varied visual data necessary to construct distinct feature spaces for every single disease category, regardless of its natural clinical rarity.

While ResNet50 is a powerful feature extractor, its reliance on localized convolutional operations restricts its ability to capture global, holistic image contexts. The future of computational dermatology lies in architectural hybridization.

- **Integration of Vision Transformers (ViTs):** Convolutional filters analyze images in small, localized patches, which can sometimes miss macro-scale diagnostic indicators (such as the overall asymmetry of a massive lesion spanning the entire image). Vision Transformers (ViTs) utilize advanced "self-attention" mechanisms that analyze the entire image in parallel, allowing the network to understand complex, long-range spatial relationships. Future frameworks will explore hybrid architectures that physically fuse the localized, high-frequency textural extraction capabilities of the ResNet50 backbone with the global contextual awareness of a Vision Transformer, yielding a drastically more comprehensive diagnostic model.
- **Ensemble Learning Paradigms:** To further boost accuracy and reduce variance, future systems will employ ensemble learning. Rather than relying on a single ResNet50 model, the diagnostic pipeline will simultaneously process the image through multiple, distinct architectures (e.g., combining ResNet50 with DenseNet121 and EfficientNet). The final diagnostic prediction will be determined via an algorithmic consensus or weighted voting mechanism among the constituent networks, significantly mitigating the individual cognitive blind spots of any single architecture.

A major barrier to the adoption of deep learning in clinical medicine is the "black box" nature of neural networks. A human dermatologist must be able to trust the machine's prediction, which is impossible if the machine cannot explain why it made that prediction.

- **Implementation of Saliency Maps and Grad-CAM:** Future iterations of this project will mandate the integration of Explainable AI (XAI) algorithms, most notably Gradient-weighted Class Activation Mapping (Grad-CAM). When the model outputs a prediction, Grad-CAM will simultaneously generate a visual, color-coded heatmap overlaying the original clinical image. This heatmap will definitively highlight the exact pixels and morphological regions that the neural network utilized to make its decision. If the model predicts "Melanoma," the human clinician can

immediately verify whether the network accurately focused on the jagged lesion border, or if it erroneously based its decision on irrelevant background noise (like a surgical skin marker). This transparency is an absolute prerequisite for regulatory approval and clinical trust.

- **Fusing Visual and Textual Metadata:** The future scope of this project envisions a transition to a multi-modal neural architecture. Alongside the ResNet50 image-processing backbone, a secondary network (such as a Recurrent Neural Network or a specialized metadata Transformer) will be constructed to simultaneously process textual patient data. By feeding the network critical metadata—including patient age, biological sex, anatomical lesion location, reported symptoms (e.g., pruritus or bleeding), and Fitzpatrick skin phototype—the AI system will be granted the necessary clinical context to disambiguate highly similar visual presentations, mirroring the holistic diagnostic process of a human physician.
- **Model Quantization for Mobile Health (mHealth):** Future development will focus heavily on model compression and quantization techniques. The deep mathematical weights of the trained ResNet50 model (typically 32-bit floating-point numbers) will be compressed into highly efficient 8-bit integers using frameworks like TensorFlow Lite or ONNX. This drastic reduction in computational size and memory footprint will allow the complex diagnostic model to be physically embedded into decentralized edge devices, such as standard consumer smartphones or portable clinical tablets.
- **Decentralized Telemedicine Portals:** By embedding the quantized model into a secure mobile application, patients in isolated, rural, or economically disadvantaged regions will gain the ability to perform immediate, off-grid preliminary screenings using their smartphone cameras. This application will be designed to instantly flag high-risk or potentially malignant lesions, automatically routing the prioritized images and localized patient data to centralized telemedicine hubs. This future integration effectively bridges the critical gap between advanced algorithmic research and tangible, life-saving clinical utility on a global scale.

REFERENCES

- 1 R. Yadav and A. Bhat, “*A systematic literature survey on skin disease detection and classification using machine learning and deep learning*,” *Multimedia Tools and Applications*, vol. 83, no. 32, pp. 78 093–78 124, 2024.
- 2 S. Inthiyaz, B. R. Altahan, S. H. Ahammad, V. Rajesh, R. R. Kalangi, L. K. Smirani, M. A. Hossain, and A. N. Z. Rashed, “*Skin disease detection using deep learning*,” *Advances in Engineering Software*, vol. 175, p. 103361, 2023.
- 3 V. Anand, S. Gupta, D. Koundal, S. R. Nayak, J. Nayak, and S. Vimal, “*Multi-class skin disease classification using transfer learning model*,” *International Journal on Artificial Intelligence Tools*, vol. 31, no. 02, p. 2250029, 2022.
- 4 O. Jaiyeoba, O. Jaiyeoba, E. Ogbuju, and F. Oladipo, “*Ai-based detection techniques for skin diseases: A review of recent methods, datasets, metrics, and challenges*,” *Journal of Future Artificial Intelligence and Technologies*, vol. 1, no. 3, pp. 318–336, 2024.
- 5 M. Sazzadul Islam Prottasha, S. Mahjabin Farin, M. Bulbul Ahmed, M. Zihadur Rahman, A. Kabir Hossain, and M. Shamim Kaiser, “*Deep learning-based skin disease detection using convolutional neural networks (cnn)*,” in *The Fourth Industrial Revolution and Beyond: Select Proceedings of IC4IR+*. Springer, 2023, pp. 551–564.
- 6 R. Alharith, A. O. Ibrahim, N. Wahid, R. Ghazali, and A. Elsafi, “*Comparative analysis of resnet models for skin cancer diagnosis: performance evaluation and insights*,” in *International Conference on Soft Computing and Data Mining*. Springer, 2024, pp. 370–378.
- 7 M. Sharma, B. Jain, C. Kargeti, V. Gupta, and D. Gupta, “*Detection and diagnosis of skin diseases using residual neural networks (resnet)*,” *International Journal of Image and Graphics*, vol. 21, no. 05, p. 2140002, 2021.
- 8 M. Elshahawy, A. Elnemr, M. Oproescu, A.-G. Schiopu, A. Elgarayhi, M. M. Elmogy, and M. Sallah, “*Early melanoma detection based on a hybrid yolov5 and resnet technique*,” *Diagnostics*, vol. 13, no. 17, p. 2804, 2023.
- 9 P. V. Raghavendra, C. Charitha, K. G. Begum, and V. S. Prasath, “*Deep learning-based skin lesion multi-class classification with global average pooling improvement*,” *Journal of Digital Imaging*, vol. 36, no. 5, pp. 2227–2248, 2023.

- 10 S. P. Choy, B. J. Kim, A. Paolino, W. R. Tan, S. M. L. Lim, J. Seo, S. P. Tan, L. Francis, T. Tsakok, M. Simpson et al., “*Systematic review of deep learning image analyses for the diagnosis and monitoring of skin disease*,” NPJ Digital Medicine, vol. 6, no. 1, p. 180, 2023.
- 11 Gairola, A. K., Kumar, V., & Sahoo, A. K. (2023). “*Deep Learning based Multiple Skin Disease Classification in Indian Territory*”. 2023 International Conference on Advancement in Computation & Computer Technologies (InCACCT), 559–564.
- 12 Mekala, P., Surendiran, B., Ramanathan, P., Ramanjaneyulu, K., Prakash, S. S. M., Chetan, T. R. V., & Varma, N. A. (2024). “*Multi-Class Skin Disease Classification: A Study of Transfer Learning Strategies for Deep Learning Models*”. 2024 International Conference on Signal Processing, Computation, Electronics, Power and Telecommunication (IconSCEPT).
- 13 Narendra, M., Harshini, T. S., & Jani Anbarasi, L. (2024). Advancing Skin Disease Diagnosis: “*A Multimodal Approach Utilizing Telegram Api Token Chatbot for Text and Image Analysis in Skin Disease Classification*”. IEEE Access, 12, 189009–189023.
- 14 Tiwari, R. G., Maheshwari, H., Gautam, V., Agarwal, A. K., & Trivedi, N. K. (2023). “*Enhancing Skin Disease Classification and Privacy Preservation through Federated Learning-Based Deep Learning*”. 2023 International Conference on Artificial Intelligence for Innovations in Healthcare Industries (ICAIIHI).

APPENDIX 1

Skin diseases Identification via Deep Learning and CNN: Using Optimized ResNet50 Architecture

Ankur Kumar

Dept. Of Computer Science and Engineering
KIET Group of Institutions
Ghaziabad, India
ankurhal22@gmail.com

Anmol Kumar

Dept. Of Computer Science and Engineering
KIET Group of Institutions
Ghaziabad, India
anmolkumar100bc@gmail.com

Anshul Rai

Dept. of Computer Science and Engineering
KIET Group of Institutions
Ghaziabad, India
raianshul33@gmail.com

Pushpendra Kumar

Dept. of Computer Science and Engineering
KIET Group of Institution
Ghaziabad, India
pushpendra.kumar@kiet.edu

Abhishek Yadav

Dept. Of Computer Science and Engineering
KIET Group of Institutions
Ghaziabad, India
abhii199works@gmail.com

Abstract—Skin disease is one of the most common human ailments but the access to the specialist dermatologist is limited. This paper present an optimized ResNet50 model for multi class skin disease classification. The main goal of this project is to improve the accuracy and generalization through a combination of transfer learning, data augmentation and regularization techniques. Automated image based classification model can provide early diagnosis and lower the risk or damage to the body. We have fine tuned a ResNet50 model, pretrained on ImageNet. The model was trained on a dataset of 22 different skin diseases. We have optimized it by extensive data augmentation techniques like random flip, rotation, zoom, contrast, a multi layer classifier and a learning rate scheduler. Training is further stabilized with early stopping. The proposed model has achieved an overall accuracy of 36.4%, weighted average F1-score of 0.31 and recall 0.27. The result shows a challenge of class imbalance in medical image classification and need for more sophisticated strategies like using class weighted and artificial data generation. Considering the size of the dataset, the optimized ResNet 50 demonstrate good performance. Future work will focus on addressing class imbalance and deploying model in clinical operations.

I. INTRODUCTION

Millions of people get affected by skin diseases every year and across all age group. [1] Conditions like acne, eczema, psoriasis and skin cancer not only affect the life of people but also pose a challenge in diagnosis due to diverse clinical presentation and overlap in visual appearance. And also access to specialist dermatologist is limited in many regions, that leads to incorrect or delayed diagnosis. An automated image based diagnosis model can help in assisting professional and provide preliminary screening. [2] Through Deep learning, par-

ticularly Convolutional Neural Network (CNN) has changed medical image analysis significantly. Models like RedNet, VCG and Inception has demonstrated remarkable performance in skin lesion classification and sometimes surpassing human expectation. Early accurate diagnosis is important, especially for skin cancer like life risking diseases. Therefore we need models like ResNet50 which is very popular because of it's balance between depth and computational efficiency. But we can not directly apply such models in real world clinical problems, there will be some problems to it. First of all medical dataset are often small, imbalanced and contain variability in lightning, scale and other imaging condition. And also the cost of acquiring a labeled dataset for rare disease lead to severe imbalance in classes of these variable disease dataset. This imbalance can make the model bias towards majority of the classes, resulting in poor generalization and unreliable prediction for uncommon diseases. To address such problems, researchers have explored various optimization strategies, including transfer learning with pre-trained weight, data augmentation, architectural modification and regularization technique etc. And also steps involved in the complete processing need to be optimized. Here we present an optimized version of ResNet50 for skin disease classification, improving the transfer learning with robust regularization. First of all fine tuned the selectively unfreeze later layer of ResNet50 pretrained on ImageNet, which helps in adaption to skin specific feature and preserving the generic visual knowledge. [3] We have also worked on regularization pipeline including data augmentation, a multi layer classifier early




Fig. 1. Pipelined for Model

stopping and learning rate scheduling. We have trained model on publicly available dataset. The result shows our model achieved moderate accuracy, higher than previous version. The only challenge that we have faced majorly here is the class imbalance.

II. LITERATURE REVIEW

One of the most popular areas in the research for Artificial Intelligence (AI) is Healthcare. From basic machine learning (ML) classifiers to more complex Deep Learning architectures, there is a huge opportunity in this field. And we can clearly see there is a significant progress in how AI is applied to various areas within healthcare systems [4]. Most of the initial AI in dermatology was based on binary classifications (i.e. distinguishing between healthy skin and lesions). But now it has become common for dermatologists to evaluate skin lesions using a multi-class severity grading system and for AI systems to identify specific lesions [5].

A. Traditional Approaches to Skin Disease Classification

Early computer diagnosis systems for skin diseases rely on hand-crafted features like color, shape descriptor, texture and dermoscopic image and feed them into classical machine learning classifiers such as support vector machine, k-nearest neighbours and random forest. These systems demonstrated feasibility but their performance was limited. These systems or models were very basic and cannot do much operations.

B. Deep Learning in Skin disease Detection

The emergence of deep convolutional neural networks marked a significant shift in medical image analysis [6]. Researchers have shown that a single CNN could achieve dermatologist-level accuracy in classifying skin lesions. But it needs a huge dataset. ResNet50 became very popular due to skip connections that allow training of very deep networks without vanishing gradients. In the recent year, EfficientNet and MobileNet have shown great performance for mobile base platforms which are lighter in size and in parameters. Other architectures such as DenseNet, Vision Transformer have also shown promising results. Deep Learning has helped a lot in optimizing the model and improving the accuracy and precision. [7] Also for other living organisms like plants we can use deep learning in monitoring their health [8].

C. ResNet50 and Transfer Learning

Transfer Learning is a technique which uses a pretrained model as a starting point for another related task. It is used in medical imaging because of pre-training on large datasets. [9] Instead of training from scratch, here we require less resource and less labeled data. [10] Through transfer learning we start working from a pretrained model which saves a lot of time and provides computational efficiency. Through ResNet50, our strategy is to freeze early layers to capture low-level features and fine-tune later layers to adapt to the target domain [11]. We can also make a hybrid model which has features of both

model and more models, this will improve the accuracy and efficiency of the model.

D. Optimization technique for CNN based Medical Image diagnosis

There are lot of factors on which the processing depends and fine tuning them will improve the performance of the model. Here for the optimization of ResNet50 we have worked on following fields:

- **Data Augmentation:** Medical datasets are sometimes small, tilted, patchy etc. So we have done some augmentation geometric transformations like: rotation, flipping, zoom, colour adjustments, and more advanced techniques like Cut and mix which increases effective sample size and improves robustness. In skin lesion classification, augmentation is very important to handle different images that vary in lighting, scale, and orientation.
- **Regularisation:** Dropout and batch normalization are standard techniques to reduce overfitting. Dropout method randomly deactivates neurons during training, forcing the network to learn redundant representations. This method help in learning new features and stop memorizing general things.
- **Learning Rate Scheduling:** Adaptive learning rate techniques help the model to converge to a better option by reducing the learning rate in case of validation loss plateaus. This method is often combined with early stopping to prevent overfitting.

III. METHODOLOGY

We have developed a structured pipeline for the proposed model focusing on increasing the accuracy and minimizing required computational resources. The overall pipeline is divided into following steps:

A. Dataset Description

We have used a publicly available skin disease dataset from KaggleHub (creatorai/skin-disease). The dataset contains dermoscopic and clinical images of various skin condition, grouped into 22 distinct disease categories [12]. For our model we have divided dataset into 80/10/10 where, 80% of the images are used for training, 10% for validation and 10% for testing. All the images are resized into 224 X 224 pixel to match the input size of ResNet50. The dataset have some class imbalance which reflect the real life condition where certain diseases are much more common than other diseases.

B. Pre-processing and Data Augmentation

All the images are pre-processed using some specific functions of ResNet50 which scales pixel values to the range expected by the pre-trained network. To increase the diversity of the training dataset and reduce overfitting, we have applied some data augmentation technique on training set such as: Random horizontal and vertical flips, Random rotation upto 20 degree, Random zoom upto 20% and Random contrast adjustment (around 10%). These augmentation techniques




Fig. 2. Dataset Distribution

were applied by TensorFlow's Sequential augmentation layer. The validation and test sets receive only the preprocessed input.

- **Normalization Layer:** It is an input preprocessing step that make ensure that the images are scaled in the same way as during the original ImageNet training.
- **Global Average Pooling:** Global Average Pooling (GAP) in ResNet50 is widely used for skin disease detection because it reduces overfitting, improves generalization, and makes the model more robust to variations in nature. [13]
- **Batch Normalization:** The normalization layer used here is Batch Normalization (BN), which is applied after each convolutional operation and before the activation function (ReLU). BN works by normalizing the output of the convolution across the mini-batch, ensuring that the activations have a consistent distribution during training.
- **Regularization:** In ResNet50, regularization is very important since the neural network is so deep that it's very likely to overfit, especially when it works with smaller, unbalanced datasets such as medical images. [14].

C. Fine-Tuning Strategy

To adapt the pre-trained ResNet50 to the dermatological domain (real world conditions) while preserving generic features, we have adopted some selective fine-tuning approach. The first 140 layers of the base model were frozen to retain low-level features learned on ImageNet, while the remaining layers are trained to capture more abstract and task specific features. The entire model was trained in a single phase as 140 layers are frozen and we only need to training remaining layers using the an optimizer with an initial learning rate of 10^{-3} , along with regularisation techniques [15].

D. Model Architecture

We have used ResNet50 model and it is pre-trained on ImagNet. The architecture is designed to use generic visual features while adapting to skin disease domain through selective fine tuning and using a custom classification as per our need. ResNet50 architecture consist of 50 layers oranigised into five main stages: an initial convolutional layer, followed by 4 residual blocks, The early layer capture low level feature

such as edges, corners and colour blobs. The remaining layer are made trainable. These deeper layer are responsible for learning more abstract, domain related features, which allows model to adapt targeted task. After base model, we create a custom classifier head to map the extracted features to the 22 skin disease classes. The head consist of Global Average Pooling which reduces feature map into a single value and the dropout layer randomly deactivate few neurons during training, forcing the network to learn different representation and thus reducing dependence on individual feature.

E. Evaluation

In this qualitative and quantitative evaluation section of the optimized ResNet50 model, we have used standard classification metrics, analyzed the confusion matrix, and examined the learning dynamics to check the effectiveness of optimization [16].

TABLE I
PER-CLASS CLASSIFICATION PERFORMANCE ON THE TEST SET.

Class Name	Precision	Recall	F1-score	Support
Acne	0.28	0.76	0.41	95
Actinic Keratosis	0.59	0.43	0.50	106
Benign Tumors	0.26	0.66	0.37	172
Bullous	0.50	0.01	0.03	76
Candidiasis	0.00	0.00	0.00	35
Drug Eruption	0.50	0.58	0.54	99
Eczema	0.31	0.64	0.42	162
Infestations/Bites	0.00	0.00	0.00	92
Lichen	0.40	0.07	0.12	85
Lupus	0.00	0.00	0.00	45
Moles	0.00	0.00	0.00	58
Psoriasis	0.39	0.28	0.33	131
Rosacea	0.00	0.00	0.00	38
Seborrheic keratoses	0.66	0.43	0.52	72
Skin Cancer	0.32	0.19	0.24	117
Sun/Sunlight Damage	0.00	0.00	0.00	53
Tinea	0.17	0.23	0.20	158
Vascular Tumors	0.00	0.00	0.00	86
Vasculitis	0.00	0.00	0.00	280
Total				1960

1) *Evaluation Metrics*: We have evaluated the model using standard classification metrics derived from the confusion matrix. For each class c (with $c = 1, 2, \dots, C$), let TP_c , FP_c , and FN_c denote the number of true positives, false positives, and false negatives, respectively. The per-class precision, recall, and F1-score are defined as:

$$\text{Precision}_c = \frac{TP_c}{TP_c + FP_c} \quad (1)$$

$$\text{Recall}_c = \frac{TP_c}{TP_c + FN_c} \quad (2)$$

$$F1_c = 2 \cdot \frac{\text{Precision}_c \cdot \text{Recall}_c}{\text{Precision}_c + \text{Recall}_c} \quad (3)$$

$$\text{Accuracy} = \frac{\sum_{c=1}^C TP_c}{N_{\text{total}}} \quad (4)$$

where N_{total} is the total number of test samples.

To summarise performance across the imbalanced set of 22 classes, we compute both macro-averaged and weighted-averaged metrics. The macro-average F1-score treats each class equally:

$$F1_{\text{macro}} = \frac{1}{C} \sum_{c=1}^C F1_c. \quad (5)$$

F1-score is harmonic mean of Precision and Recall. The weighted-average F1-score weights each class by its support i.e. number of true samples:

$$F1_{\text{weighted}} = \frac{\sum_{c=1}^C TP_c \cdot F1_c}{\sum_{c=1}^C TP_c}. \quad (6)$$

2) *Confusion Matrix*: The confusion matrix provide a detailed view of misclassification. Many samples from minority classes are incorrectly assigned to few high support classes which makes the majority classes dominating in prediction. Some clinically similar conditions are confusing as they reflect the visual overlap of scaly lesions. And some rows shows no true positives, their prediction are spread across other classes. Thee confusion matrix highlights the problem of the class imbalance.




Fig. 3. Confusion Matrix

3) *Loss Curves and Accuracy Graphs*: The training and validation accuracy and loss curves over 50 epochs training period shows that the training accuracy increases steadily while validation accuracy plateaus early, here the widening gap suggest that there is a mild overfitting. The loss in validation stabilises after 10 epochs approximately and the learning rate scheduler shows there is a reduction in learning rate when validation loss stops improving. Early stopping prevent further overfitting by stoping training so that best model is restored.

4) Hardware Configuration:

- **Processor**: Intel Core i7 / AMD Ryzen 7 (Equivalent or higher).




Fig. 4. Accuracy and Loss Graph

- **GPU:** NVIDIA T4 GPU (via Google Colab Pro) for high-speed tensor computations.
- **VRAM:** 16GB High-Bandwidth Memory (HBM).
- **System RAM:** 12.7GB (Standard Colab Environment).

5) *Software Environment:*

- **Operating System:** Linux-based cloud environment.
- **Programming Language:** Python 3.10.
- **Deep Learning Framework:** TensorFlow 2.x and Keras API.
- **Data Analysis Libraries:** NumPy, Pandas, and Matplotlib.
- **Documentation Tools:** LaTeX code and Overleaf for IEEE formatted conference paper.

IV. RESULTS

In the result section we will cover all the parameter and factors involved in the calculating the result. It represent the quantitative and qualitative evaluation of the optimised ResNet50 model.

A. Overall Performance

The model has achieved a test accuracy of 36.4%. The macro-average precision, recall, and F1-score were each 0.27, while the weighted-average F1-score was 0.31. These values are significantly higher than random guessing, showcasing that the model has learned meaningful patterns. However, the low macro-averages indicate severe difficulty in classifying minority classes.

B. Per-Class Metrics

The table in the methodology section represents precision, recall, and F1-score for each of the 22 classes, along with test set support. Classes with zero recall (e.g., Candidiasis, Infestations/Bites, Lupus, Moles, Rosacea, Sun/Sunlight Damage, Vascular Tumors, Vasculitis) were never correctly identified.

And classes with larger support, such as Benign tumors (recall 0.66) and Eczema (recall 0.64), achieved moderate recall but suffered low precision, indicating that many predictions were false positives.

C. Confusion Matrix Analysis

- **Majority-class bias:** The model tends to predict a small set of high-support classes (e.g., Benign tumors, Eczema, Tinea) for many test samples, causing rows of minority classes to be spread across these columns.
- **Clinically similar confusions:** Tinea (class 16) is frequently misclassified as Eczema (class 6) and Seborrheic keratoses (class 13), reflecting visual overlap of scaly lesions. Actinic keratosis (class 1) and Sun/Sunlight damage (class 15) also show mutual confusion.
- **Zero-performance rows:** Several classes have no true positive predictions, meaning the model fails entirely to recognise those conditions which has lesser size of the dataset.

D. Training Dynamics

The accuracy-loss graph shows that the training and validation accuracy/loss over 50 epochs. Training accuracy steadily increases, while validation accuracy plateaus after approximately 10 epochs. The widening gap indicates mild overfitting, which is lowered by early stopping and learning rate reduction. The validation loss is stabilized, and the best model is restored from the epoch with the lowest validation loss.

V. CONCLUSION

In this paper, we present an optimised ResNet50 model for multi-class skin disease classification. The optimisation strategy combined transfer learning technique with selective fine-tuning, data augmentation, regularisation classifier head with dropout, early stopping, and learning rate scheduling. [17]

The model was evaluated on a real-world dataset of 22 skin disease categories exhibiting severe class imbalance.




Fig. 5. Output Images

Our results shows that the optimised model which achieves an overall accuracy of 36.4%, with a macro-average F1-score of 0.27. While the model performs reasonably on a few majority classes, it fails entirely on many minority conditions. The confusion matrix reveals that misclassifications are accumulated among similar diseases and that the model's predictions are biased toward high-support classes. These findings highlights the critical challenge of class imbalance in medical image classification and demonstrate that even with advanced optimisation techniques, standard fine-tuning is insufficient when rare diseases have less representation. The main contributions of this work are:

- A systematic fine-tuning and regularisation pipeline for ResNet50 tailored to skin disease classification.
- A detailed analysis of per-class performance and confusion patterns, providing insights for future research.

VI. LIMITATIONS

There are some limitation in this project specifically in the training the model and dataset used which leads to have some huge impact in the result section like: The dataset used here was not perfectly balanced and no explicit class balancing mechanism was used. This leads to poor performance in case of minority classes. The model was trained on clinical images with varying condition so generalisation to other dataset may be limited.

VII. FUTURE WORK

Future work will focus on improving the performance of the proposed model and making it easy to use and also available to all of us. But the major work should be done particularly for

minority classes and making this model, through the following directions:

- Manage cost-sensitive learning or focus on increasing dataset of rare diseases for better training and accuracy.
- Explore other models so that we can optimise current one by extracting important features from other models and improving the performance.
- Evaluating the model on external datasets to assess generalisability. We can also find defect or mistake in our current model by working on external datasets.
- Integrate explainability methods such as Grad-CAM to provide justification for the predictions, which is essential for clinical acceptance.
- Deploy the model as a mobile-based screening tool for preliminary diagnosis in resource-constrained environment.

REFERENCES

- [1] R. Yadav and A. Bhat, "A systematic literature survey on skin disease detection and classification using machine learning and deep learning," *Multimedia Tools and Applications*, vol. 83, no. 32, pp. 78 093–78 124, 2024.
- [2] S. Inthiyaz, B. R. Altahan, S. H. Ahammad, V. Rajesh, R. R. Kalangi, L. K. Smirani, M. A. Hossain, and A. N. Z. Rashed, "Skin disease detection using deep learning," *Advances in Engineering Software*, vol. 175, p. 103361, 2023.
- [3] V. Anand, S. Gupta, D. Koundal, S. R. Nayak, J. Nayak, and S. Vimal, "Multi-class skin disease classification using transfer learning model," *International Journal on Artificial Intelligence Tools*, vol. 31, no. 02, p. 2250029, 2022.
- [4] O. N. Pise and S. Borde, "Skin disease detection using machine learning," *Int. J. Eng. Res. Technol.(IJERT)*, vol. 9, pp. 64–68, 2020.
- [5] O. Jaiyeoba, O. Jaiyeoba, E. Ogbuju, and F. Oladipo, "Ai-based detection techniques for skin diseases: A review of recent methods, datasets, metrics, and challenges," *Journal of Future Artificial Intelligence and Technologies*, vol. 1, no. 3, pp. 318–336, 2024.
- [6] M. Sazzadul Islam Prottasha, S. Mahjabin Farin, M. Bulbul Ahmed, M. Zihadur Rahman, A. Kabir Hossain, and M. Shamim Kaiser, "Deep learning-based skin disease detection using convolutional neural networks (cnn)," in *The Fourth Industrial Revolution and Beyond: Select Proceedings of IC4IR+*. Springer, 2023, pp. 551–564.
- [7] P. Kumar, S. Arora, and S. Arora, "Real-time sugarcane bud detection using efficient lightweight convolutional blocks," *Arabian Journal for Science and Engineering*, 2025. [Online]. Available: <https://doi.org/10.1007/s13369-025-10906-3>
- [8] —, "Real time sugarcane stem health detection with yolov8 tinier and repncsp elan," *Scientific Reports*, vol. 15, no. 1, p. 43672, 2025.
- [9] R. Alharith, A. O. Ibrahim, N. Wahid, R. Ghazali, and A. Elsafi, "Comparative analysis of resnet models for skin cancer diagnosis: performance evaluation and insights," in *International Conference on Soft Computing and Data Mining*. Springer, 2024, pp. 370–378.
- [10] M. Sharma, B. Jain, C. Kargeti, V. Gupta, and D. Gupta, "Detection and diagnosis of skin diseases using residual neural networks (resnet)," *International Journal of Image and Graphics*, vol. 21, no. 05, p. 2140002, 2021.
- [11] M. Elshahawy, A. Elnemr, M. Oproescu, A.-G. Schiopu, A. Elgarayhi, M. M. Elmogy, and M. Sallah, "Early melanoma detection based on a hybrid yolov5 and resnet technique," *Diagnostics*, vol. 13, no. 17, p. 2804, 2023.
- [12] S. P. Choy, B. J. Kim, A. Paolino, W. R. Tan, S. M. L. Lim, J. Seo, S. P. Tan, L. Francis, T. Tsakok, M. Simpson *et al.*, "Systematic review of deep learning image analyses for the diagnosis and monitoring of skin disease," *NPI Digital Medicine*, vol. 6, no. 1, p. 180, 2023.
- [13] P. V. Raghavendra, C. Charitha, K. G. Begum, and V. S. Prasath, "Deep learning-based skin lesion multi-class classification with global average pooling improvement," *Journal of Digital Imaging*, vol. 36, no. 5, pp. 2227–2248, 2023.

Proof of Acceptance



Ankur Kumar <ankurhal22@gmail.com>

Acceptance Notification – PTEMS 2026

1 message

Microsoft CMT <noreply@microsoft.com>
To: Ankur Kumar <ankurhal22@gmail.com>

Wed, Mar 18, 2026 at 10:00 AM

Dear Ankur Kumar
Greetings from PTEMS 2026.

Congratulations!

With reference to your research paper having Paper ID: 458
titled: "Skin diseases Identification via Deep Learning and CNN: Using Optimized ResNet50 Architecture" submitted to the
INTERNATIONAL CONFERENCE ON PROGRESSIVE TRENDS IN ENGINEERING, MANAGEMENT & SCIENCE (PTEMS-2026),

We are pleased to inform you that your submission has been ACCEPTED for oral presentation and publication in PTEMS-2026,
subject to the incorporation of the reviewer's comments and compliance with conference policies.

Registration & Camera-Ready Submission Guidelines